



## EECS 3221: Operating System Fundamentals

Hamzeh Khazaei  
Department of Electrical Engineering and  
Computer Science

Week 10:  
**Virtual Memory**

# Assignment Project Exam Help



March 15, 2021  
EECS3221: Operating System Fundamentals  
10.1  
Virtual Memory

1

<https://powcoder.com>



## Add WeChat powcoder Chapter 10: Virtual Memory

---

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



EECS3221: Operating System Fundamentals  
10.2  
Virtual Memory

2



## Background

- Code needs to be in memory to execute, **but entire program rarely used**
  - Error code, unusual routines, large data structures
- Entire program code not needed at same time
- Consider ability to execute partially-loaded program
  1. Program no longer constrained by limits of physical memory
  2. Each program takes less memory while running -> more programs run at the same time
    - Increased CPU utilization and throughput with no increase in response time or turnaround time
  3. Less I/O needed to load or swap programs into memory -> each user program runs faster

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.3

Virtual Memory

3

<https://powcoder.com>



## Add WeChat powcoder

### Virtual Memory

- **Virtual memory** – separation of user logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation
  - More programs running concurrently
  - Less I/O needed to load or swap processes



EECS3221: Operating System Fundamentals

10.4

Virtual Memory

4

2



## Virtual memory (Cont.)

- **Virtual address space** – logical view of how process is stored in memory
  - Usually start at address 0, contiguous addresses until end of space
  - Meanwhile, physical memory organized in page frames
  - MMU must map logical to physical
  
- Virtual memory can be implemented via:
  - Demand paging
  - Demand segmentation

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.5

Virtual Memory

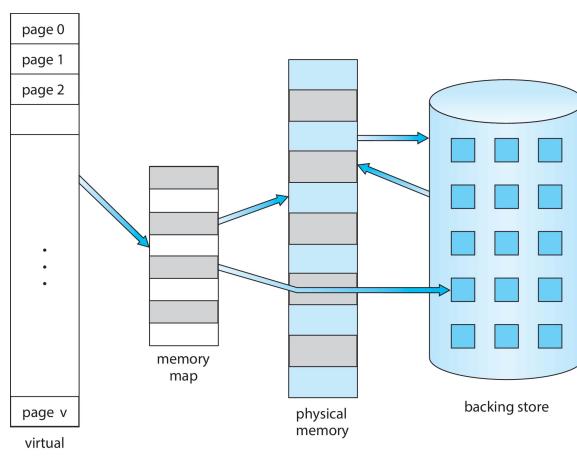
5

<https://powcoder.com>



## Add WeChat powcoder

Virtual Memory That is Larger Than Physical Memory



EECS3221: Operating System Fundamentals

10.6

Virtual Memory

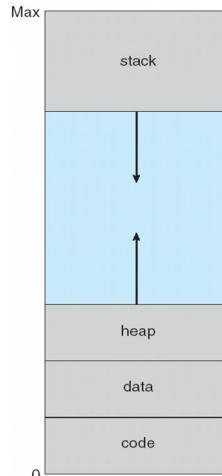
6

3



## Virtual-address Space

- Usually design logical address space for stack to start at Max logical address and grow “down” while heap grows “up”
  - Maximizes address space use
  - Unused address space between the two is hole
    - ▶ No physical memory needed until heap or stack grows to a given new page
- Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc
- System libraries shared via mapping into virtual address space
- Shared memory by mapping pages read-write into virtual address space
- Pages can be shared during `fork()`, speeding process creation



## Assignment Project Exam Help

EECS3221: Operating System Fundamentals

10.7

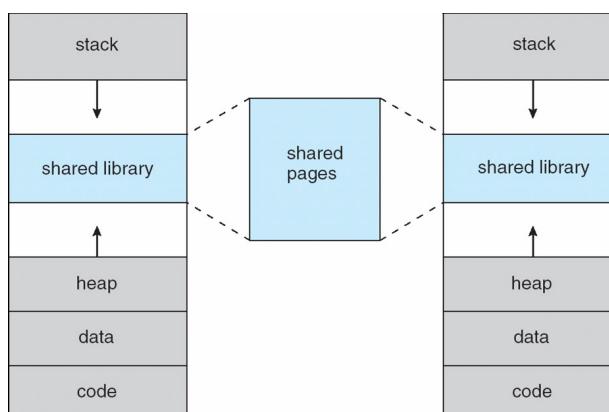
Virtual Memory

7

<https://powcoder.com>



## Add WeChat powcoder Shared Library Using Virtual Memory



EECS3221: Operating System Fundamentals

10.8

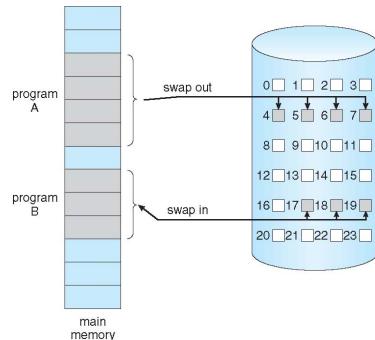
Virtual Memory

8



## Demand Paging

- Could bring entire process into memory at load time
- Or bring a page into memory only when it is needed
  - Less I/O needed, no unnecessary I/O
  - Less memory needed
  - Faster response
  - More users
- Like paging system with swapping (diagram on right)
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
- **Lazy swapper** – never swaps a page into memory unless page will be needed
  - Swapper that deals with pages is a **pager**



## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.9

Virtual Memory

9

<https://powcoder.com>



## Add WeChat powcoder

### Valid-Invalid Bit

- With each page table entry, a valid-invalid bit is associated (**v**  $\Rightarrow$  in-memory – **memory resident**, **i**  $\Rightarrow$  not-in-memory)
- Initially valid-invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table

- During MMU address translation, if valid-invalid bit in page table entry is **i**  $\Rightarrow$  **page fault**

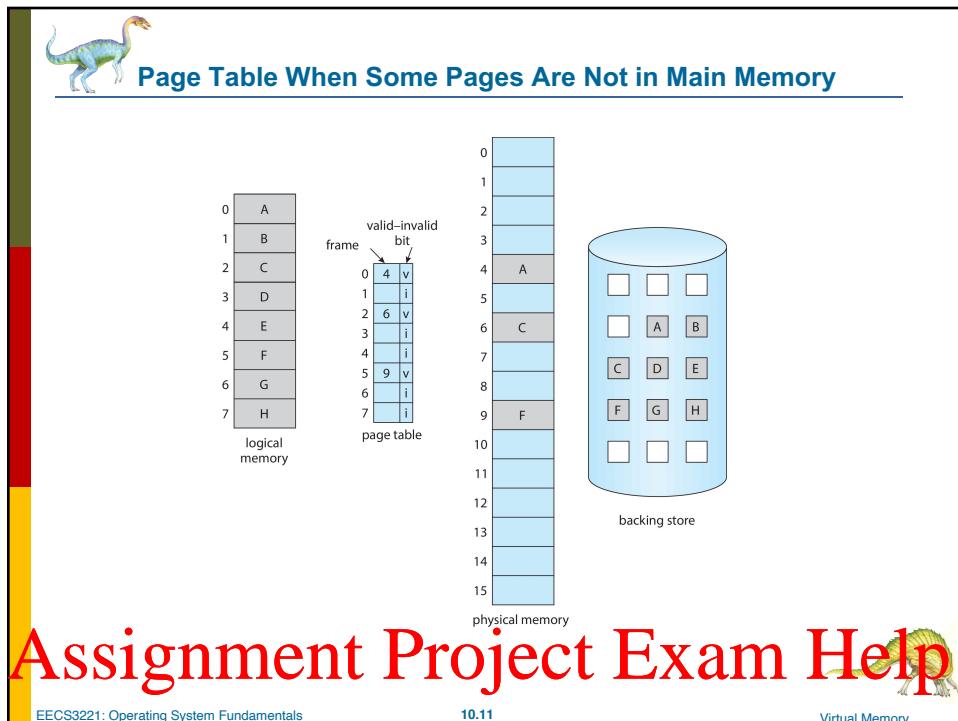


EECS3221: Operating System Fundamentals

10.10

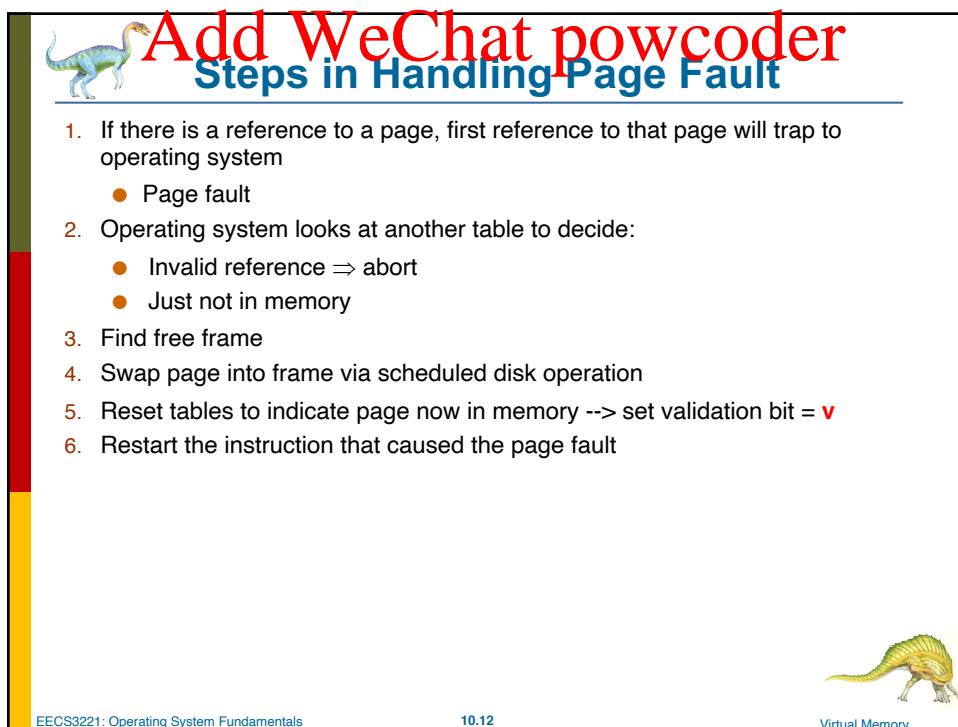
Virtual Memory

10

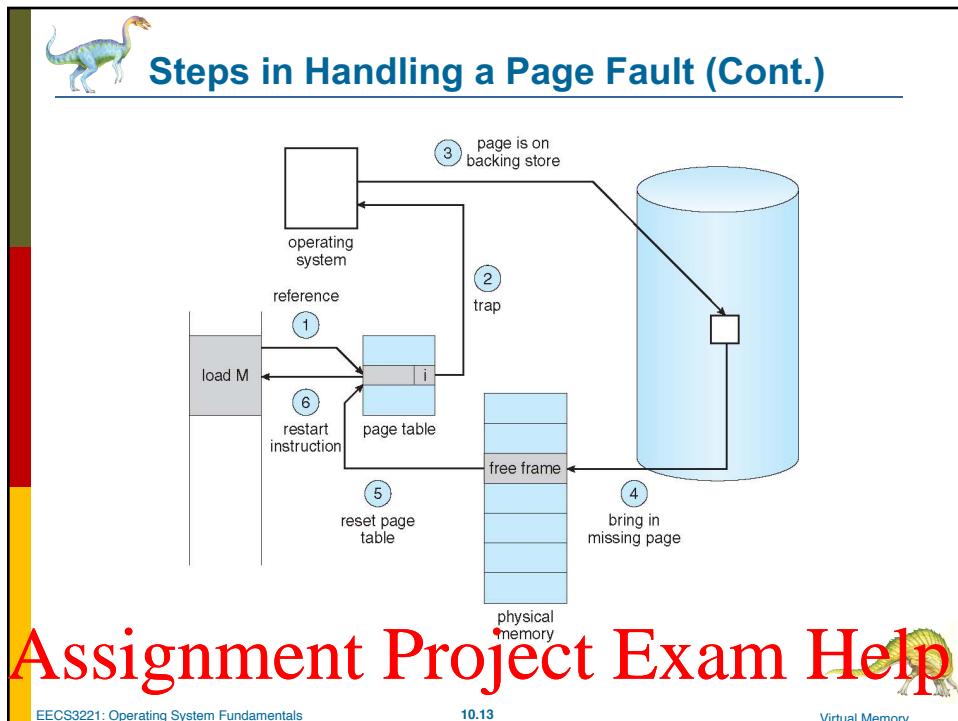


11

<https://powcoder.com>



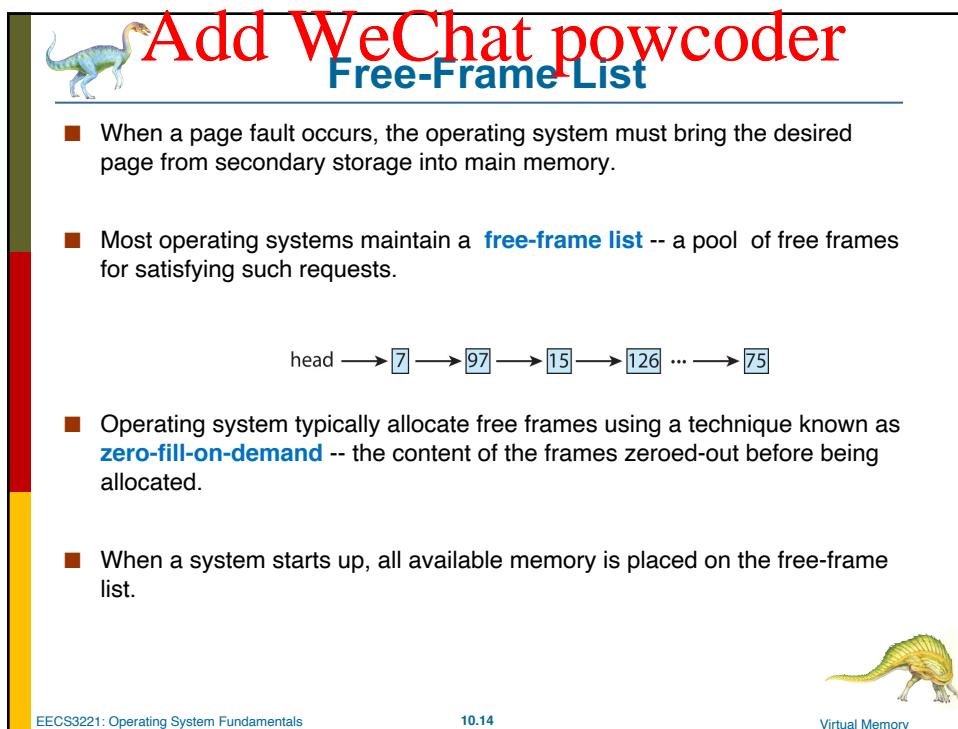
12



EECS3221: Operating System Fundamentals      10.13      Virtual Memory

13

<https://powcoder.com>



EECS3221: Operating System Fundamentals

10.14

Virtual Memory

14



## Performance of Demand Paging

- Three major activities
  - Service the interrupt – careful coding means just several hundred instructions needed
  - Read the page – lots of time
  - Restart the process – again just a small amount of time
- Page Fault Rate  $0 \leq p \leq 1$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault
- Effective Access Time (EAT)

$$EAT = (1 - p) \times \text{memory access} + p (\text{page fault overhead} + \text{swap page out} + \text{swap page in})$$

## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.15

Virtual Memory

15

<https://powcoder.com>



## Add WeChat powcoder Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$ 
$$= ((1 - p) \times 200) + (p \times 8,000,000)$$
$$= 200 + (p \times 7,999,800)$$
- If one access out of 1,000 causes a page fault ( $p=0.001$ ), then  
 $EAT = 8.2 \text{ microseconds}$ .  
This is a slowdown by a factor of 40!!
- If want performance degradation < 10 percent
  - $220 > 200 + 7,999,800 \times p$   
 $20 > 7,999,800 \times p$
  - $p < 0.0000025$
  - $p <$  one page fault in every 400,000 memory accesses



EECS3221: Operating System Fundamentals

10.16

Virtual Memory

16



## Copy-on-Write

- **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
  - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
  - Designed to have child call `exec()`
  - Very efficient

# Assignment Project Exam Help

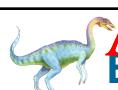
EECS3221: Operating System Fundamentals

10.17

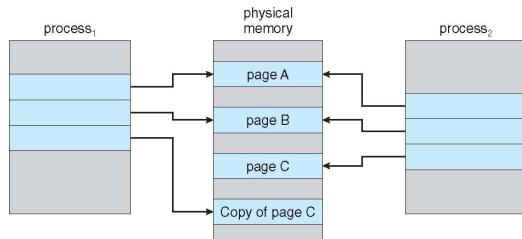
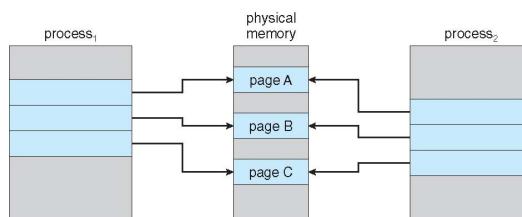
Virtual Memory

17

<https://powcoder.com>



## Add WeChat powcoder Before Process 1 Modifies Page C



EECS3221: Operating System Fundamentals

10.18

Virtual Memory

18



## What Happens if There is no Free Frame?

- Used up by process pages
- Also, in demand from the kernel, I/O buffers, etc.
- Algorithm – terminate? swap out? replace the page?
  - Page replacement – find some page in memory, but not really in use, page it out
- Performance – want an algorithm which will result in minimum number of page faults
- How much to allocate to each?
- Same page may be brought into memory several times

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.19

Virtual Memory

19

<https://powcoder.com>



## Add WeChat powcoder

### Page Replacement

- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory



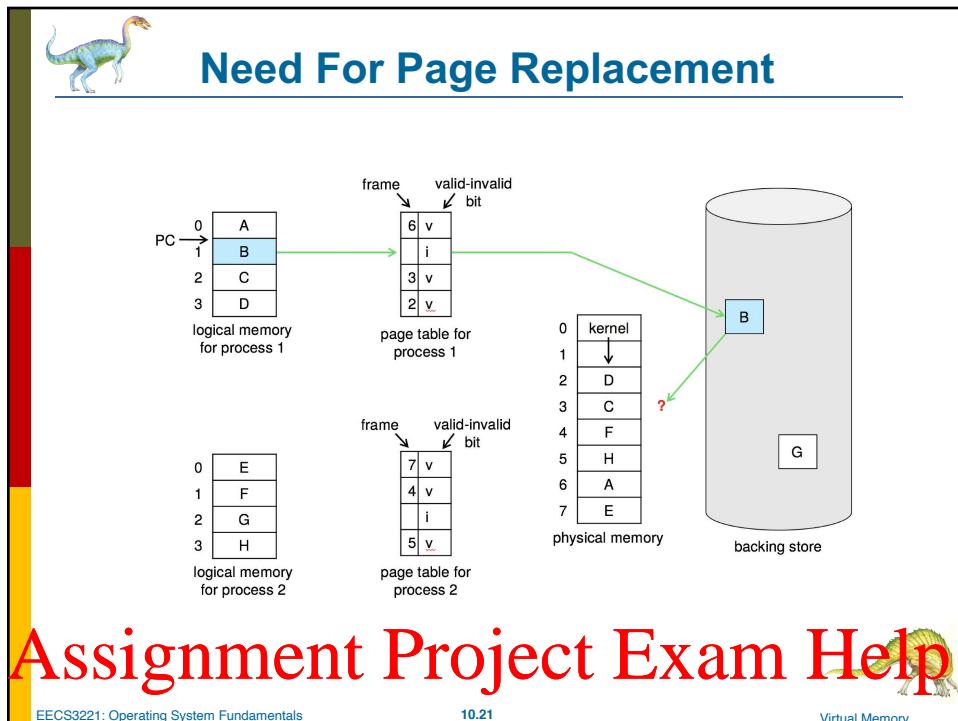
EECS3221: Operating System Fundamentals

10.20

Virtual Memory

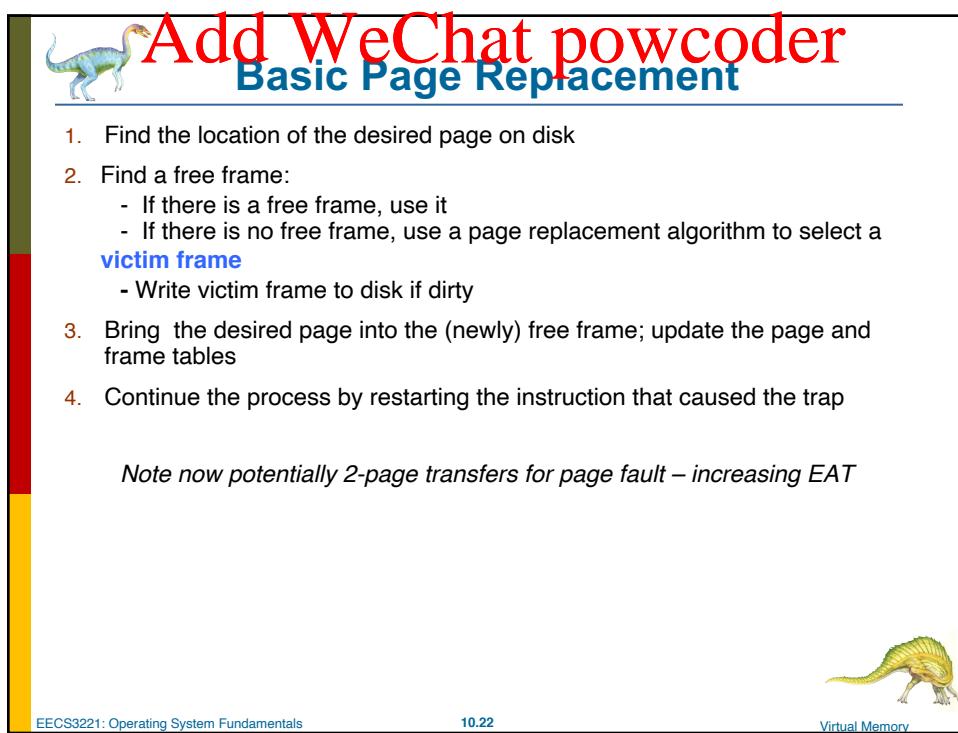
20

10



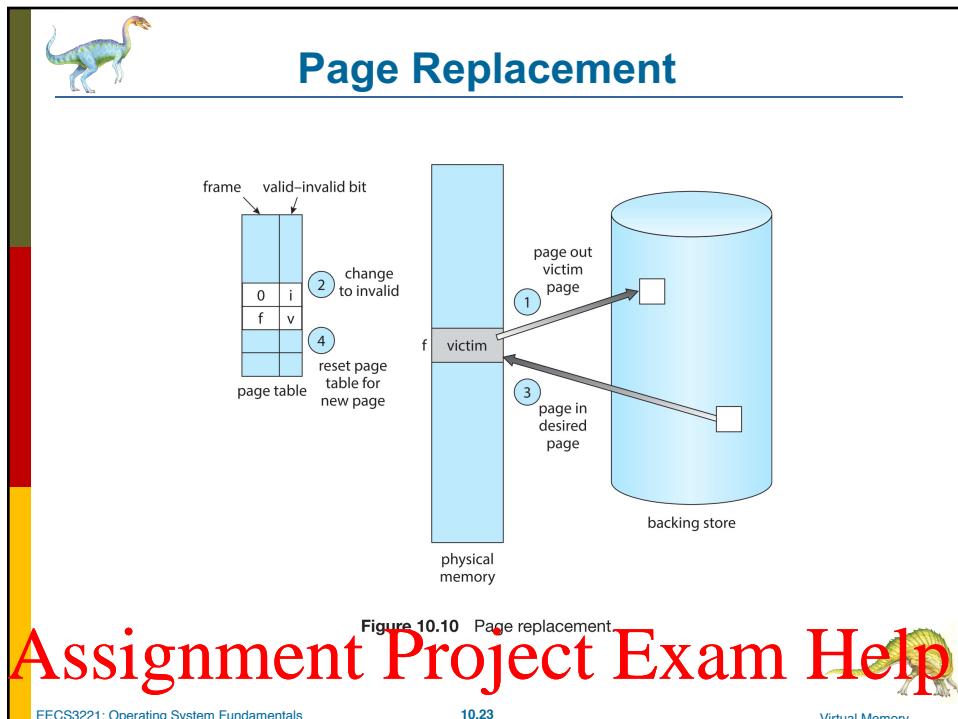
EECS3221: Operating System Fundamentals      10.21      Virtual Memory

21 <https://powcoder.com>



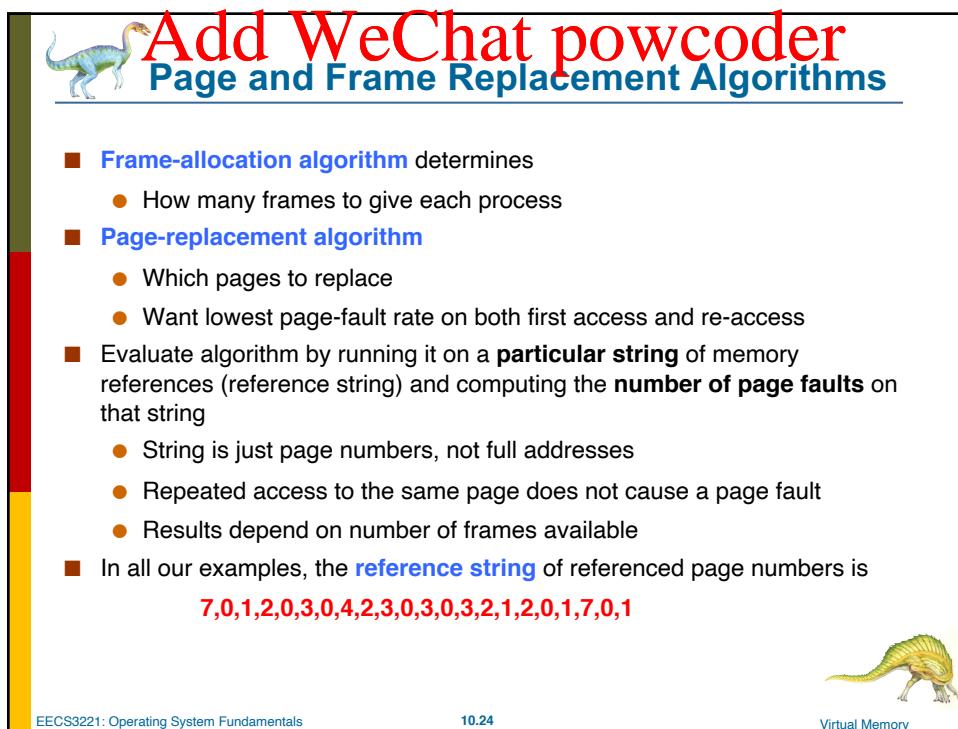
EECS3221: Operating System Fundamentals      10.22      Virtual Memory

22

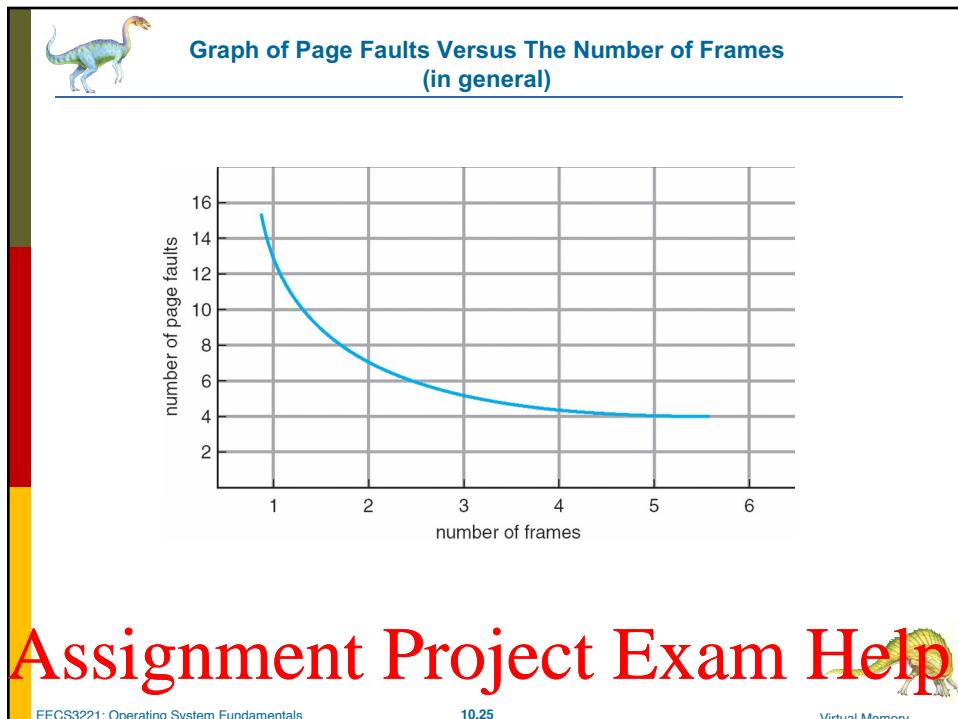


23

<https://powcoder.com>



24



## Assignment Project Exam Help

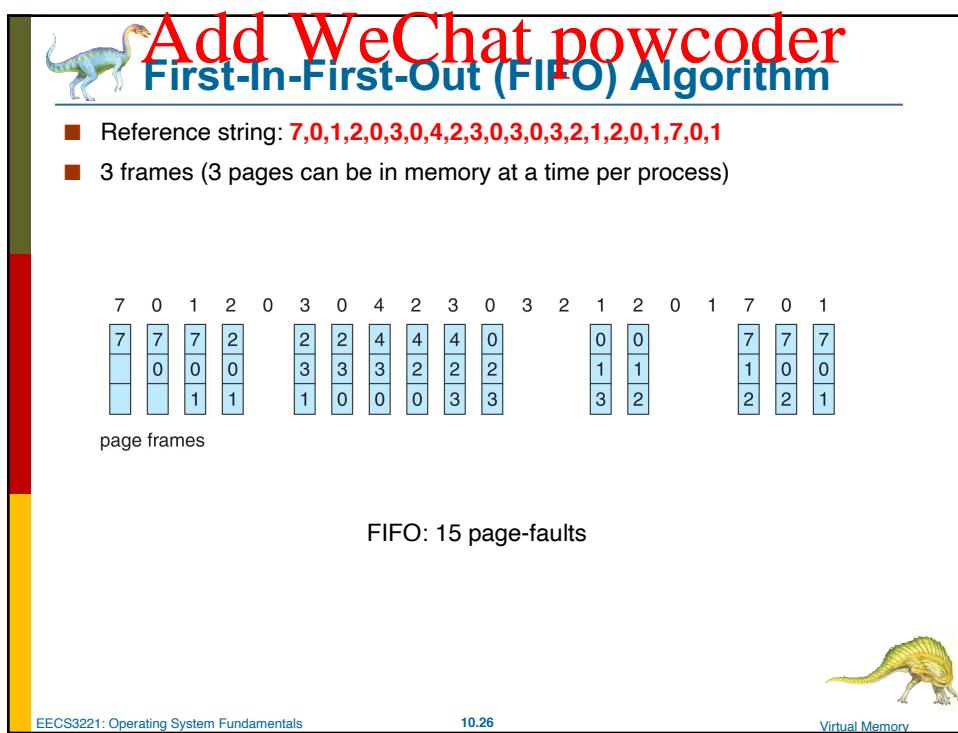
EECS3221: Operating System Fundamentals

10.25

Virtual Memory

25

<https://powcoder.com>



EECS3221: Operating System Fundamentals

10.26

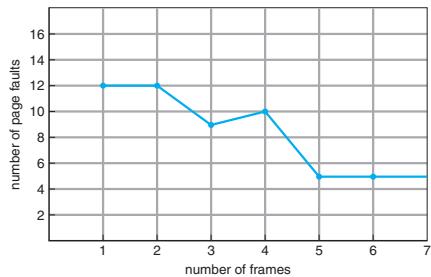
Virtual Memory

26



## FIFO Illustrating Belady's Anomaly

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - Adding more frames can cause more page faults!
- ▶ Belady's Anomaly



- How to track ages of pages?
- Just use a FIFO queue

## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.27

Virtual Memory

27

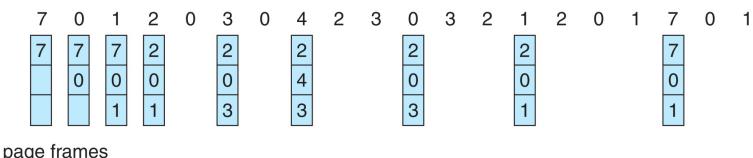
<https://powcoder.com>



## Add WeChat powcoder

### Optimal Algorithm

- Replace page that will not be used for longest period of time
  - 9 page-faults is optimal for the example
- How do you know this?
  - Can't read the future
- Used for measuring how well your algorithm performs



page frames

- 9 faults – absolute best.



EECS3221: Operating System Fundamentals

10.28

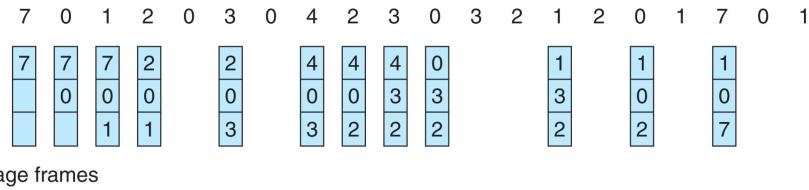
Virtual Memory

28



## Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page



- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.29

Virtual Memory

29

<https://powcoder.com>



## Add WeChat powcoder LRU Algorithm (Cont.)

- Counter implementation
  - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
  - When a page needs to be changed, look at the counters to find smallest value
    - ▶ Search through table needed and
    - ▶ A write to memory for each memory access
- Stack implementation
  - Keep a stack of page numbers in a double link form:
  - Page referenced:
    - ▶ move it to the top
    - ▶ requires 6 pointers to be changed
  - But each update more expensive
  - The tail pointer points to the bottom of the stack, which is the LRU page
  - No search for replacement



EECS3221: Operating System Fundamentals

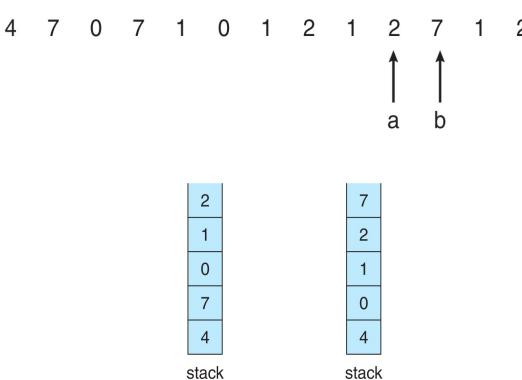
10.30

Virtual Memory

30

 Use Of A Stack to Record Most Recent Page References

LRU (like OPT) is a case of **stack algorithms** that doesn't have Belady's Anomaly.



EECS3221: Operating System Fundamentals      10.31      Virtual Memory

31

<https://powcoder.com>

 Add WeChat powcoder  
LRU Approximation Algorithms

- LRU needs special hardware and still slow
- **Reference bit**
  - With each page associate a bit, initially = 0
  - When page is referenced, bit set to 1
  - Replace any with reference bit = 0 (if one exists)
    - ▶ We do not know the order, however
- **Second-chance algorithm**
  - Generally, FIFO, plus hardware-provided reference bit
  - **Clock** replacement
  - If page to be replaced has
    - ▶ Reference bit = 0 -> replace it
    - ▶ reference bit = 1 then:
      - set reference bit 0, leave page in memory
      - replace next page, subject to same rules

EECS3221: Operating System Fundamentals      10.32      Virtual Memory

32

EECS3221: Operating System Fundamentals      10.33      Virtual Memory

33

<https://powcoder.com>

EECS3221: Operating System Fundamentals

10.34

Virtual Memory

34



## Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- **Least Frequently Used (LFU) Algorithm** replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used
- Not common – potential problems?
  
- One solution is to shift the counts right by 1 bit at regular intervals, forming an exponentially decaying average usage count.

## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.35

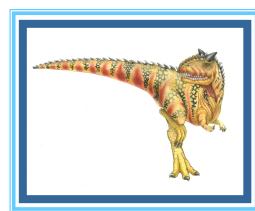
Virtual Memory

35

<https://powcoder.com>

Add WeChat powcoder

## Second Module



Operating System Concepts – 10<sup>th</sup> Edition

Silberschatz, Galvin and Gagne ©2018

36

18



## Add-on procedures – best practices

- Other procedures are often used in addition to a specific page-replacement algorithm:
  - Keep a pool of free frames, always
    - ▶ Then frame available when needed, not to be found at fault time
    - ▶ Read page into free frame and select victim to evict and add to free pool
    - ▶ When convenient, evict victim
  - Possibly, keep list of modified pages
    - ▶ When backing store idle, write pages there and set to non-dirty
  - Possibly, keep free frame contents intact and note what is in them
    - ▶ If referenced again before reused, no need to load contents again from disk
    - ▶ Generally useful to reduce penalty if wrong victim frame selected

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.37

Virtual Memory

37

<https://powcoder.com>



## Add WeChat powcoder Applications and Page Replacement

- All these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e., databases
- Memory intensive applications can cause double buffering
  - OS keeps copy of page in memory as I/O buffer
  - Application keeps page in memory for its own work
- Operating system can give direct access to the disk, getting out of the way of the applications
  - **Raw disk** mode
  - Bypasses buffering, locking, etc.



EECS3221: Operating System Fundamentals

10.38

Virtual Memory

38

19



## Allocation of Frames

- Each process needs **minimum** number of frames
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
  - instruction is 6 bytes, might span 2 pages
  - 2 pages to handle *from*
  - 2 pages to handle *to*
- **Maximum** of course is total frames in the system
- Two major allocation schemes
  - fixed allocation
  - priority allocation
- Many variations

## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.39

Virtual Memory

39

<https://powcoder.com>



## Add WeChat powcoder

### Fixed Allocation

- Equal allocation – For example, if there are 93 frames (after allocating frames for the OS) and 5 processes, give each process 18 frames
  - Keep 3 frames as free frame buffer pool
- Proportional allocation – Allocate according to the size of process
  - Dynamic as degree of multiprogramming, process sizes change

–  $s_i$  = size of process  $p_i$

$$m = 64$$

–  $S = \sum s_i$

$$s_1 = 10$$

–  $m$  = total number of frames

$$s_2 = 127$$

–  $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$$a_1 = \frac{10}{137} \times 62 \approx 4$$

$$a_2 = \frac{127}{137} \times 62 \approx 57$$

EECS3221: Operating System Fundamentals

10.40

Virtual Memory

40





## Global vs. Local Allocation

- Now that we know allocation, let's get back to page replacement:

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
  - But then process execution time can vary greatly
  - But greater throughput so **more common**
- **Local replacement** – each process selects from only its own set of allocated frames
  - More consistent per-process performance
  - But possibly underutilized memory

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.41

Virtual Memory

41

<https://powcoder.com>



## Add WeChat powcoder Priority Allocation

- Use a proportional allocation scheme but using **priorities** rather than size
- If process  $P_i$  generates a page fault,
  - select for replacement one of its frames
  - select for replacement a frame from a process with lower priority number

EECS3221: Operating System Fundamentals

10.42

Virtual Memory

42





## Reclaiming Pages

- A strategy to implement global page-replacement policy
- All memory requests are satisfied from the free-frame list
  - The free-frame list will not be lower/upper than a threshold
- Page replacement is triggered when the list falls below a certain threshold.
- This strategy attempts to ensure there is always enough free memory to satisfy new requests.

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.43

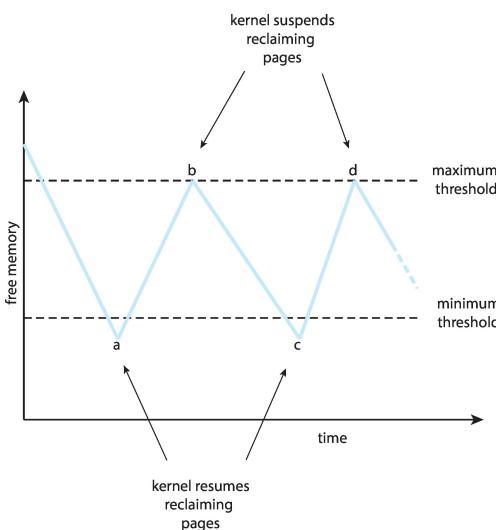
Virtual Memory

43

<https://powcoder.com>



## Add WeChat powcoder Reclaiming Pages Example (reaper procedure)



EECS3221: Operating System Fundamentals

10.44

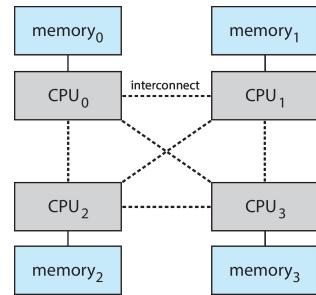
Virtual Memory

44



## Non-Uniform Memory Access

- So far, all memory accessed equally
- Many systems are **NUMA** – speed of access to memory varies
  - Consider system boards containing CPUs and memory, interconnected over a system bus
- NUMA multiprocessing architecture



## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.45

Virtual Memory

45

<https://powcoder.com>



## Add WeChat powcoder

### Non-Uniform Memory Access (Cont.)

- Optimal performance comes from allocating memory “close to” the CPU on which the thread is scheduled
  - And modifying the scheduler to schedule the thread on the same system board when possible
  - Solved by Solaris by creating **Igroups**
    - ▶ Structure to track CPU / Memory low latency groups
    - ▶ Used my schedule and pager
    - ▶ When possible, schedule all threads of a process and allocate all memory for that process within the Igroup



EECS3221: Operating System Fundamentals

10.46

Virtual Memory

46



## Thrashing

- A process is thrashing if it is spending more time paging than executing.
- If a process does not have “enough” pages, the page-fault rate is very high
  - Page fault to get page
  - Replace existing frame
  - But quickly need replaced frame back
- This leads to:
  - ▶ Low CPU utilization
  - ▶ Operating system thinking that it needs to increase the degree of multiprogramming
  - ▶ Another process added to the system – the problem might get exacerbated.

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.47

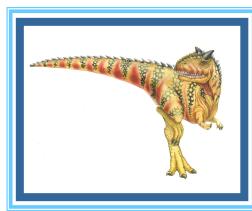
Virtual Memory

47

<https://powcoder.com>

Add WeChat powcoder

## Third Module



Operating System Concepts – 10<sup>th</sup> Edition

Silberschatz, Galvin and Gagne ©2018

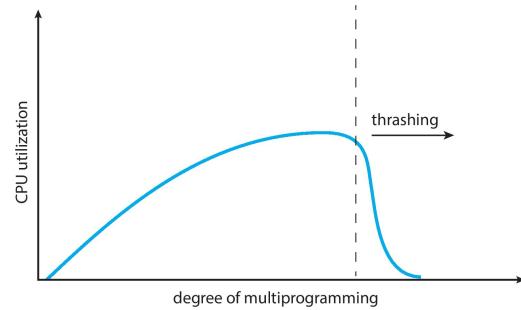
48

24



## Thrashing (Cont.)

- **Thrashing.** A process is busy swapping pages in and out



# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.49

Virtual Memory

49

<https://powcoder.com>



## Add WeChat powcoder Demand Paging and Thrashing

- Why does demand paging work?
  - Locality model
    - Process migrates from one locality to another
    - Localities may overlap
- Why does thrashing occur?
$$\Sigma \text{ size of locality} > \text{total memory size}$$
- We can limit the effect by using local or priority page replacement

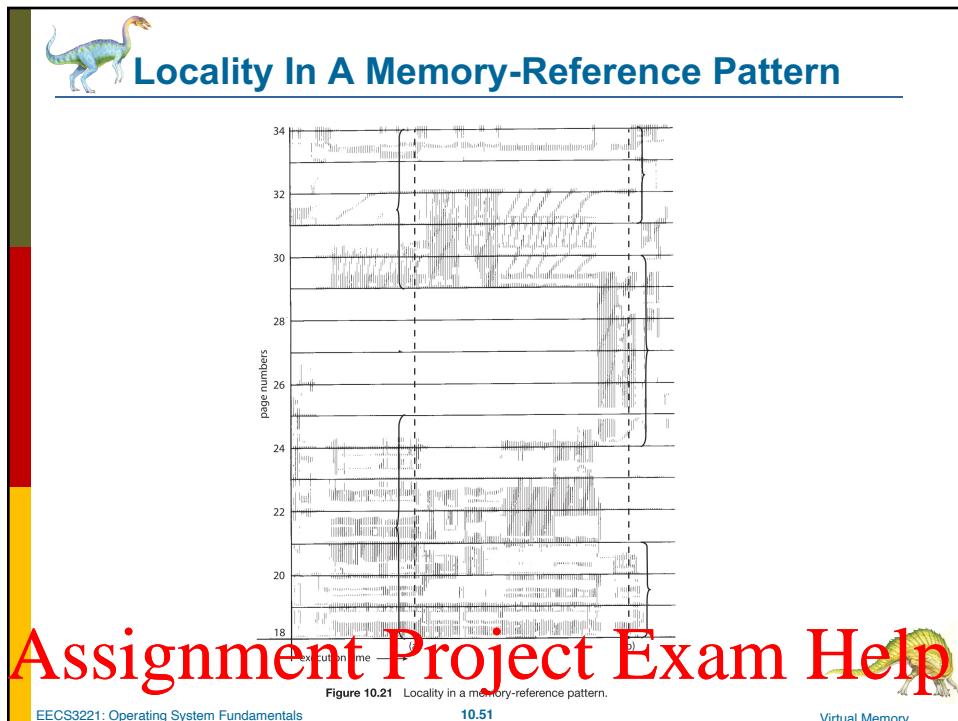


EECS3221: Operating System Fundamentals

10.50

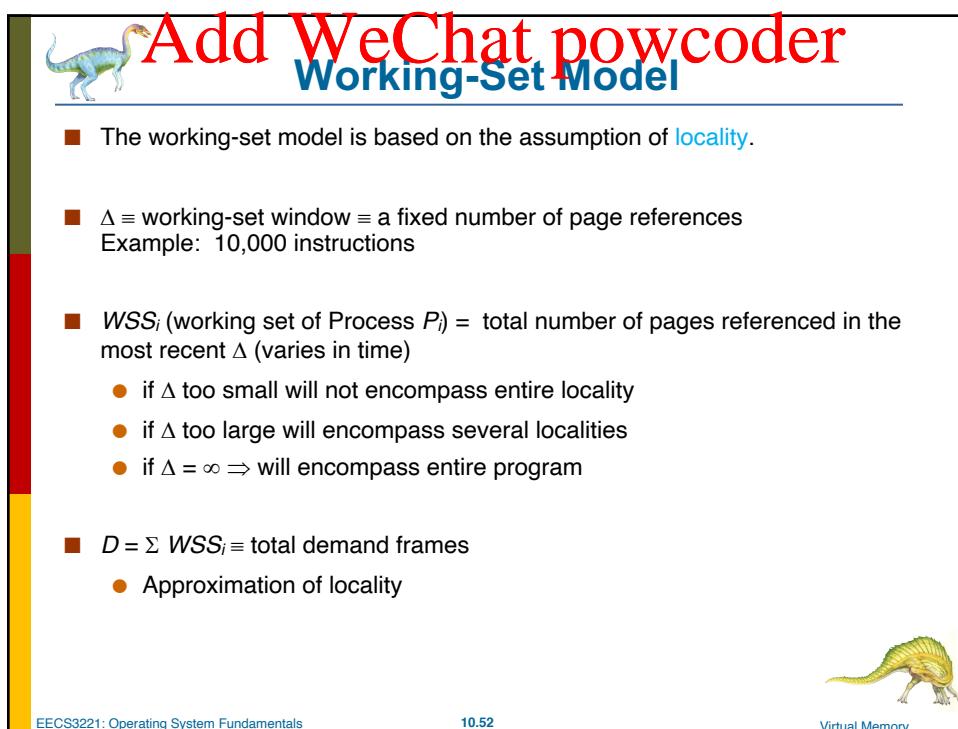
Virtual Memory

50



51

<https://powcoder.com>

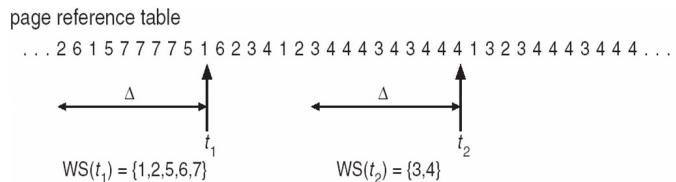


52



## Working-Set Model (Cont.)

- if  $D > m \Rightarrow$  Thrashing
- Policy if  $D > m$ , then suspend or swap out one of the processes
- Following picture, delta = 10



## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.53

Virtual Memory

53

<https://powcoder.com>



## Add WeChat powcoder Keeping Track of the Working Set

- Approximate with interval timer + a reference bit
- Example:  $\Delta = 10,000$ 
  - Timer interrupts after every 5000 time-units
  - Keep in memory 2 bits for each page
  - Whenever a timer interrupt occurs, copy and sets the values of all reference bits to 0
  - If one of the bits in memory = 1  $\Rightarrow$  page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time-units
  - Any problems?



EECS3221: Operating System Fundamentals

10.54

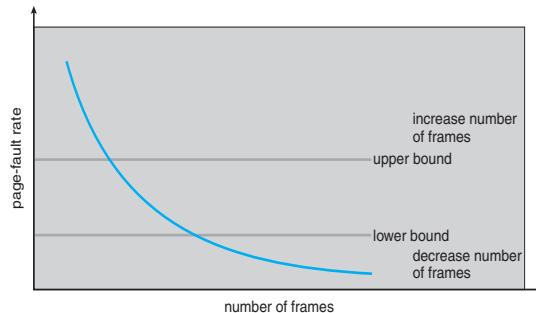
Virtual Memory

54



## Page-Fault Frequency

- More direct approach than WSS
- Establish “acceptable” **page-fault frequency (PFF)** rate and use **local replacement policy**
  - If actual rate too low, process loses frame
  - If actual rate too high, process gains frame



## Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.55

Virtual Memory

55

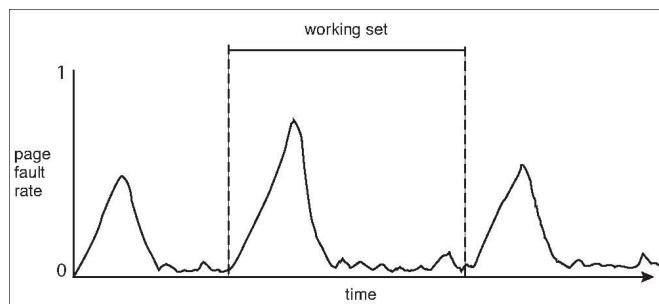
<https://powcoder.com>



## Add WeChat powcoder

### Working Sets and Page Fault Rates

- Direct relationship between working set of a process and its page-fault rate
- Working set changes over time
- Peaks and valleys over time



EECS3221: Operating System Fundamentals

10.56

Virtual Memory

56



## Allocating Kernel Memory

- Treated differently from user memory
  - Allocating and deallocating memory frequently
- Often allocated from a free-memory pool
  - Kernel requests memory for structures of varying sizes
  - Some kernel memory needs to be contiguous
    - ▶ i.e., for device I/O
- Now, we examine two strategies for managing free memory that is assigned to kernel processes:
  - The “buddy system” and
  - The “slab allocation”

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.57

Virtual Memory

57

<https://powcoder.com>



## Add WeChat powcoder Buddy System

- Allocates memory from fixed-size segment consisting of physically-contiguous pages
- Memory allocated using **power-of-2 allocator**
  - Satisfies requests in units sized as power of 2
  - Request rounded up to next highest power of 2
  - When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
    - ▶ Continue until appropriate sized-chunk become available
- For example, assume 256KB chunk available, kernel requests 21KB
  - Split into  $A_L$  and  $A_R$  of 128KB each
    - ▶ One further divided into  $B_L$  and  $B_R$  of 64KB
      - One further into  $C_L$  and  $C_R$  of 32KB each – one used to satisfy request
- Advantage – quickly **coalesce** unused chunks into larger chunk
- Disadvantage – ?
  - internal fragmentation
    - ▶ 50% wastage of memory



EECS3221: Operating System Fundamentals

10.58

Virtual Memory

58

**Buddy System Allocator**

physically contiguous pages

```

graph TD
    A[256 KB] --> B[128 KB AL]
    A --> C[128 KB AR]
    B --> D[64 KB BL]
    B --> E[64 KB BR]
    D --> F[32 KB CL]
    D --> G[32 KB CR]
    
```

The diagram illustrates the Buddy System Allocator's memory hierarchy. It starts with a 256 KB block at the top, which is split into two 128 KB blocks, labeled  $A_L$  and  $A_R$ . Each of these 128 KB blocks is further split into two 64 KB blocks, labeled  $B_L$  and  $B_R$ . Finally, each of these 64 KB blocks is split into two 32 KB blocks, labeled  $C_L$  and  $C_R$ . All blocks are represented by rectangles with black outlines, and they are arranged in a tree-like structure where each node has two children.

**Assignment Project Exam Help**

EECS3221: Operating System Fundamentals      10.59      Virtual Memory

59

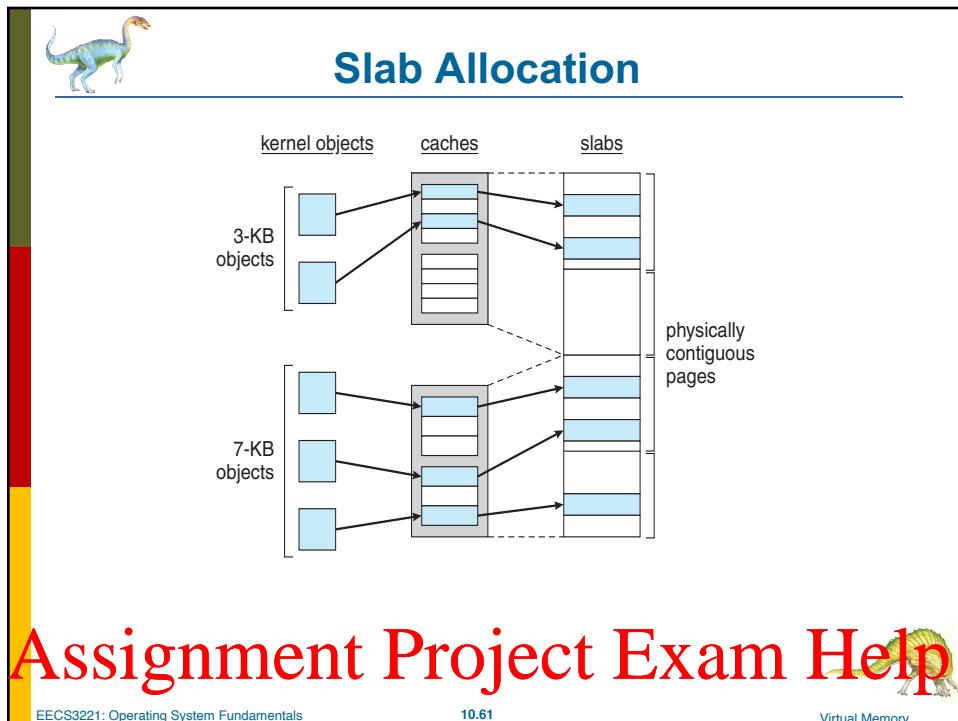
<https://powcoder.com>

**Add WeChat powcoder Slab Allocator**

- Alternate strategy
- **Slab** is one or more physically contiguous pages
- **Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
  - Each cache filled with **objects** – instantiations of the data structure
- When cache created, filled with objects marked as **free**
- When structures stored, objects marked as **used**
- If slab is full of used objects, next object allocated from empty slab
  - If no empty slabs, new slab allocated
- Benefits include no fragmentation, fast memory request satisfaction

EECS3221: Operating System Fundamentals      10.60      Virtual Memory

60



## Assignment Project Exam Help

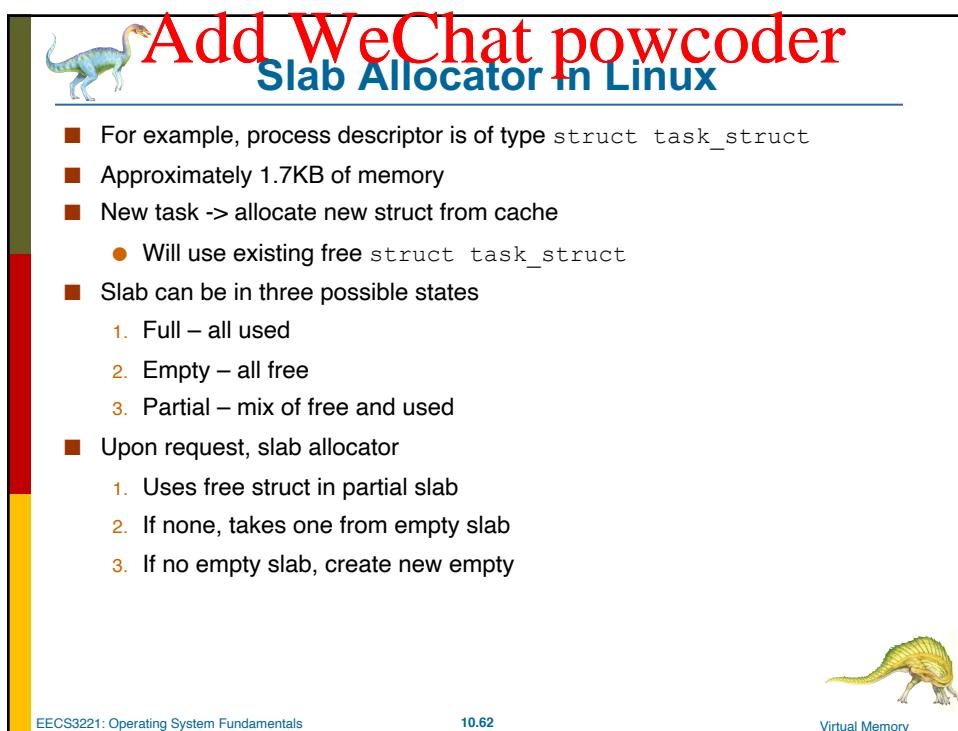
EECS3221: Operating System Fundamentals

10.61

Virtual Memory

61

<https://powcoder.com>



EECS3221: Operating System Fundamentals

10.62

Virtual Memory

62



## Slab Allocator in Linux (Cont.)

- Slab started in Solaris, now wide-spread for both kernel mode and user memory in various OSes
- Linux 2.2 had SLAB, now has both SLOB and SLUB allocators
  - SLOB for systems with limited memory, like embedded systems
    - ▶ Simple List of Blocks – maintains 3 list objects for small (256 byte), medium (less than 1024 byte), large objects
    - ▶ Memory requests are allocated from an object on the appropriate list using a first-fit policy.
  - SLUB, beginning at version 2.6.24, is performance-optimized SLAB
    - ▶ removes per-CPU queues (memory saving in multi-cpu systems)
    - ▶ metadata stored in page structure

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.63

Virtual Memory

63

<https://powcoder.com>



## Add WeChat powcoder

### Other Considerations

- Prepaging
- Page size
- TLB reach
- Program structure
- I/O interlock and page locking



EECS3221: Operating System Fundamentals

10.64

Virtual Memory

64



## Prepaging

- To reduce the large number of page faults that occurs at process startup
- Pre-page all or some of the pages a process will need, before they are referenced
- But if pre-paged pages are unused, I/O and memory was wasted
- Assume  $s$  pages are pre-paged and  $a\%$  of the pages is used
  - Is cost of  $s * a$  save pages faults > or < than the cost of pre-paging  $s * (100 - a)$  unnecessary pages?
  - $a$  near zero  $\Rightarrow$  pre-paging loses

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.65

Virtual Memory

65

<https://powcoder.com>



## Add WeChat powcoder Page Size

- Sometimes OS designers have a choice
  - Especially if running on custom-built CPU
- Page size selection must take into consideration:
  - Fragmentation
  - Page table size
  - Resolution
  - I/O overhead
  - Number of page faults
  - Locality
  - TLB size and effectiveness
- Always power of 2, usually in the range  $2^{12}$  (4,096 bytes) to  $2^{22}$  (4,194,304 bytes)
- On average, growing over time



EECS3221: Operating System Fundamentals

10.66

Virtual Memory

66



## TLB Reach

- TLB Reach - The amount of memory accessible from the TLB
- $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- Ideally, the working set of each process is stored in the TLB
  - Otherwise, there is a high degree of TLB miss
- Increase the Page Size
  - This may lead to an increase in fragmentation as not all applications require a large page size
- Provide Multiple Page Sizes
  - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.67

Virtual Memory

67

<https://powcoder.com>



## Add WeChat powcoder Program Structure

- Program structure
  - `int[128,128] data;`
  - Each row is stored in one page of size  $2^{12}$
  - Program 1

```
for (j = 0; j < 128; j++)
    for (i = 0; i < 128; i++)
        data[i,j] = 0;
```

$128 \times 128 = 16,384$  page faults

- Program 2

```
for (i = 0; i < 128; i++)
    for (j = 0; j < 128; j++)
        data[i,j] = 0;
```

128 page faults



EECS3221: Operating System Fundamentals

10.68

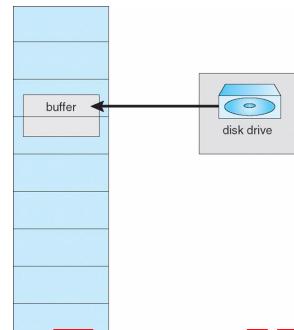
Virtual Memory

68



## I/O interlock

- **I/O Interlock** – Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm
- **Pinning** of pages to lock into memory



# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.69

Virtual Memory

69

<https://powcoder.com>



## Add WeChat powcoder Operating System Examples

- Linux
- Solaris
- Windows



EECS3221: Operating System Fundamentals

10.70

Virtual Memory

70



## Windows

- Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page
- Processes are assigned **working set minimum** and **working set maximum**
- Working set minimum is the minimum number of pages the process is guaranteed to have in memory
- A process may be assigned as many pages as possible up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum

# Assignment Project Exam Help



EECS3221: Operating System Fundamentals

10.71

Virtual Memory

71

<https://powcoder.com>



## Add WeChat powcoder Performance of Demand Paging

- Stages in Demand Paging (worse case)
  1. Trap to the operating system
  2. Save the user registers and process state
  3. Determine that the interrupt was a page fault
  4. Check that the page reference was legal and determine the location of the page on the disk
  5. Issue a read from the disk to a free frame:
    1. Wait in a queue for this device until the read request is serviced
    2. Wait for the device seek and/or latency time
    3. Begin the transfer of the page to a free frame
  6. While waiting, allocate the CPU to some other user/program
  7. Receive an interrupt from the disk I/O subsystem (I/O completed)
  8. Save the registers and process state for the other user/program
  9. Determine that the interrupt was from the disk
  10. Correct the page table and other tables to show page is now in memory
  11. Wait for the CPU to be allocated to this process again (process now in ready queue)
  12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction



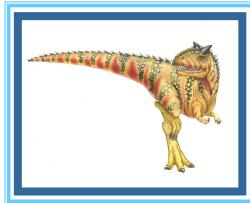
EECS3221: Operating System Fundamentals

10.72

Virtual Memory

72

**End of Chapter 10.**



**Assignment Project Exam Help**

Operating System Concepts – 10<sup>th</sup> Edition

Silberschatz, Galvin and Gagne ©2018

73

<https://powcoder.com>

Add WeChat powcoder