



## EECS 3221: OPERATING SYSTEM FUNDAMENTALS

Hamzeh Khazaei  
Department of Electrical Engineering and  
Computer Science

Week 9, Module 1:  
**Main Memory**

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.1

Main Memory

1

<https://powcoder.com>



## Add WeChat powcoder Chapter 9: Memory Management

- Background
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Swapping
- Example: The Intel 32 and 64-bit Architectures
- Example: ARMv8 Architecture

EECS3221: Operating System Fundamentals

9.2

Main Memory

2



## Background

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of:
  - addresses + read requests, or
  - address + data and write requests
- Register access is done in one CPU clock
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.3

Main Memory

3


<https://powcoder.com>



## Add WeChat powcoder

### Protection

- Need to ensure that a process can access only those addresses in its address space.
- We can provide this protection by using a pair of **base** and **limit registers** to define the logical address space of a process



EECS3221: Operating System Fundamentals

9.4

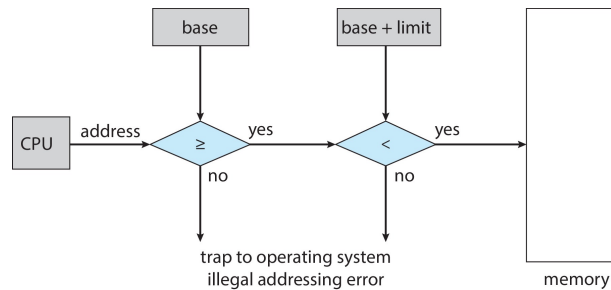
Main Memory

4



## Hardware Address Protection

- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



- The instructions to loading the base and limit registers are **privileged**

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.5

Main Memory

5

<https://powcoder.com>



## Add WeChat powcoder

### Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
- Addresses represented in different ways at different stages of a program's life
  - Source code addresses usually symbolic (int x)
  - Compiled code addresses **bind** to relocatable addresses
    - ▶ i.e. "14 bytes from beginning of this module"
  - Linker or loader will bind relocatable addresses to absolute addresses
    - ▶ i.e. 74014
  - Each binding maps one address space to another

EECS3221: Operating System Fundamentals

9.6

Main Memory

6



## Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
  - **Compile time:** If starting point of memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
  - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
  - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
    - ▶ Need hardware support for address maps (e.g., base and limit registers)
    - ▶ Most OSs use this method. Let's see how this can be done?

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.7

Main Memory


7

<https://powcoder.com>



## Add WeChat powcoder

### Multistep Processing of a User Program



EECS3221: Operating System Fundamentals

9.8

Main Memory

8



## Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
- logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.9

Main Memory


9

<https://powcoder.com>



## Memory-Management Unit (MMU)

- Hardware device that at **run time** maps virtual to physical address



- Many methods possible, covered in the rest of this chapter

EECS3221: Operating System Fundamentals

9.10


Main Memory

10



## Memory-Management Unit (Cont.)

- Consider simple scheme, which is a generalization of the base-register scheme.
- The base register now called **relocation register**
- The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
  - Execution-time binding occurs when reference is made to location in memory
  - Logical address bound to physical addresses



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.11

Main Memory

11

<https://powcoder.com>



## Add WeChat powcoder

### Dynamic Loading

- The entire program does need to be in memory to execute
- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - Implemented through program design
  - OS can help by providing libraries to implement dynamic loading

EECS3221: Operating System Fundamentals

9.12

Main Memory

12



## Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the binary program image
- Dynamic linking – linking postponed until execution time
- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
  - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**
- Consider applicability to patching system libraries
  - Versioning may be needed
- Unlike dynamic loading, dynamic linking and shared libraries generally require help from the operating system

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.13

Main Memory

13

<https://powcoder.com>



## Add WeChat powcoder

### Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is **one early method**
- Main memory usually is divided into **two partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes then held in high memory
  - Each process contained in single contiguous section of memory

EECS3221: Operating System Fundamentals

9.14

Main Memory

14



## Contiguous Allocation (Cont.)

- **Relocation registers** used to protect user processes from each other, and from changing operating-system code and data
  - Base register contains value of smallest physical address
  - Limit register contains range of logical addresses – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*
  - Can then allow actions such as kernel code being **transient** and kernel changing size

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.15

Main Memory

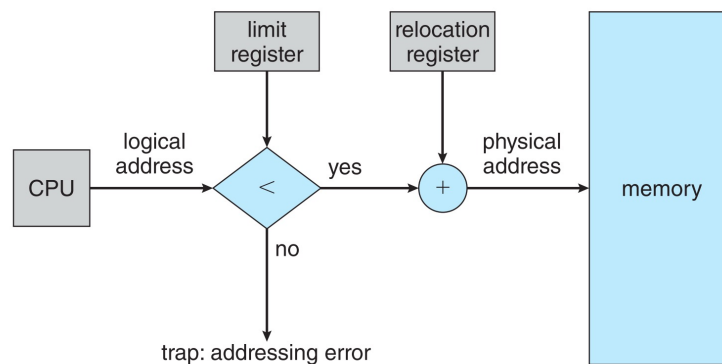
15

<https://powcoder.com>



## Add WeChat powcoder

### Hardware Support for Relocation and Limit Registers



EECS3221: Operating System Fundamentals

9.16

Main Memory

16






## Variable Partition

### ■ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
  - a) allocated partitions
  - b) free partitions (hole)



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.17

Main Memory

17

<https://powcoder.com>



## Add WeChat powcoder Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- **First-fit**: Allocate the **first** hole that is big enough
- **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit**: Allocate the **largest** hole; must also search entire list
  - Produces the largest leftover hole

*First-fit and best-fit better than worst-fit in terms of **speed** and **storage utilization**.*

*First fit is **faster** than best fit but there is **no difference** in terms of **storage utilization**.*

EECS3221: Operating System Fundamentals

9.18

Main Memory

18



## Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given  $N$  blocks allocated,  $N/2$  blocks lost to fragmentation
  - 1/3 may be unusable -> **50-percent rule**

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.19

Main Memory

19

<https://powcoder.com>



## Add WeChat powcoder

### Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time (not compilation or load time)
  - I/O problem
    - ▶ Latch job in memory while it is involved in I/O
    - ▶ Do I/O only into OS buffers
- Fragmentation is a general problem in computing that can occur wherever we must manage blocks of data. (disk, external storages, etc.)
- Another possible solution to the external-fragmentation problem is to permit the physical address space of processes to be **noncontiguous**

EECS3221: Operating System Fundamentals

9.20

Main Memory

20



## Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size **N** pages, need to find **N** free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.21

Main Memory

21

<https://powcoder.com>



## Add WeChat powcoder

### Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** (**p**) – used as an index into a **page table** which contains base address of each page in physical memory
  - **Page offset** (**d**) – combined with base address to define the physical memory address that is sent to the memory unit

page number	page offset
p	d
m - n	n

- For given logical address space  $2^m$  and page size  $2^n$

EECS3221: Operating System Fundamentals

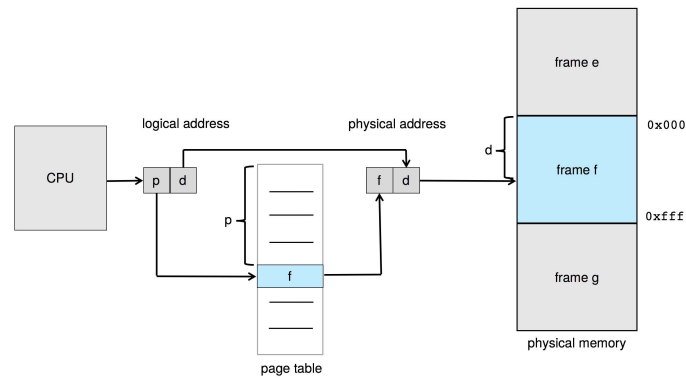
9.22

Main Memory

22



## Paging Hardware



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.23

Main Memory


23

<https://powcoder.com>




## Add WeChat powcoder

### Paging Model of Logical and Physical Memory




logical  
memory



page table

frame  
number



physical  
memory

EECS3221: Operating System Fundamentals

9.24


Main Memory

24



## Paging Example

- Logical address:  $n = 2$  and  $m = 4$ . Using a page size of 4 bytes and a physical memory of 32 bytes (8 pages)



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.25


Main Memory

25

<https://powcoder.com>

Add WeChat powcoder

## Second Module



26



## Paging -- Calculating internal fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of 2,048 - 1,086 = 962 bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = 1 / 2 frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track
- Page sizes growing over time
  - Solaris supports two page sizes – 8 KB and 4 MB
  - In Linux and Mac run:

```
bash:> getconf PAGESIZE
```

# Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.27

Main Memory


27

<https://powcoder.com>



## Add WeChat powcoder

### Free Frames



Before allocation

After allocation

EECS3221: Operating System Fundamentals

9.28

Main Memory

28



## Implementation of Page Table

- Page table is kept in main memory
  - **Page-table base register (PTBR)** points to the page table
  - **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the actual data or instruction
  - This is expensive; any idea how we can make this better?
- The two-memory access problem can be solved using a special fast-lookup hardware cache called **translation look-aside buffers (TLBs)** (also called **associative memory**).

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.29

Main Memory

29

<https://powcoder.com>



## Add WeChat powcoder

### Translation Look-Aside Buffer

- Each entry in the TLB consists of two parts: a key (or tag) and a value.
- When the associative memory is presented with an item, the item is compared with all keys simultaneously.
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access
  - Typically, TLB entries for key kernel code are wired down.

EECS3221: Operating System Fundamentals

9.30

Main Memory

30



## Hardware

### ■ Associative memory – parallel search

Page #	Frame #

### ■ Address translation (p, d)

- If p is in associative register, get frame # out
- Otherwise get frame # from page table in memory

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.31

Main Memory


31

<https://powcoder.com>



## Add WeChat powcoder

### Paging Hardware With TLB



EECS3221: Operating System Fundamentals

9.32

Main Memory

32





## Effective Access Time

- Hit ratio – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
  - If we find the desired page in TLB then a mapped-memory access take 10 ns
  - Otherwise, we need two memory access, so it is 20 ns
- **Effective Access Time (EAT)**  
$$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$
implying 20% slowdown in access time
- Consider amore realistic hit ratio of 99%,  
$$\text{EAT} = 0.99 \times 10 + 0.01 \times 20 = 10.1\text{ns}$$
implying only 1% slowdown in access time.

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.33

Main Memory

33

<https://powcoder.com>



## Memory Protection in Page Table

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel

EECS3221: Operating System Fundamentals

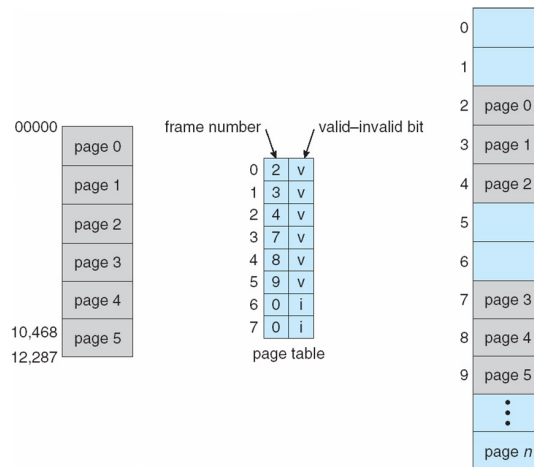
9.34

Main Memory

34



## Valid (v) or Invalid (i) Bit In A Page Table



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.35

Main Memory

35

<https://powcoder.com>



## Add WeChat powcoder

### Shared Pages

#### ■ Shared code

- One copy of read-only code shared among processes (i.e., text editors, compilers, window systems)
- Like multiple threads sharing the same process space
- Also useful for inter-process communication if sharing of read-write pages is allowed

#### ■ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

EECS3221: Operating System Fundamentals

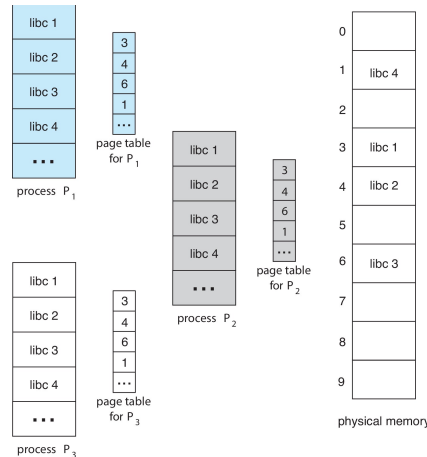
9.36

Main Memory

36



## Shared Pages Example



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.37

Main Memory

37

<https://powcoder.com>



## Add WeChat powcoder

### Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ )
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes → each process 4 MB of physical address space for the page table alone
    - Don't want to allocate that contiguously in main memory
    - Want to divide the page table into smaller pieces

page number	page offset
p	d

m - n                      n

- One simple solution is to divide the page table into smaller units
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables

EECS3221: Operating System Fundamentals

9.38


Main Memory

38



## Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.39

Main Memory

39

<https://powcoder.com>



## Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
  - a page number consisting of 20 bits
  - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
  - a 10-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
10	10	12

- Where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

EECS3221: Operating System Fundamentals


9.40

Main Memory

40



## Address-Translation Scheme



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.41

Main Memory

41

<https://powcoder.com>



## 64-bit Logical Address Space

- Even two-level paging scheme not enough
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes

EECS3221: Operating System Fundamentals

9.42

Main Memory

42



## Three-level Paging Scheme

- One solution is to add a 2<sup>nd</sup> outer page table
- But in the following example the 2<sup>nd</sup> outer page table is still 2<sup>34</sup> bytes in size (i.e., 16 GB)
  - And possibly 4 memory access to get to one physical memory location

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12

- You can see from this example why, for 64-bit architectures, hierarchical page tables are generally considered inappropriate.

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.43

Main Memory

43

<https://powcoder.com>



## Add WeChat powcoder

### Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted

EECS3221: Operating System Fundamentals


9.44

Main Memory

44



## Hashed Page Table



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.45

Main Memory

45

<https://powcoder.com>



## Add WeChat powcoder

### Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, one page table to track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs

EECS3221: Operating System Fundamentals



9.46

Main Memory

46



## Inverted Page Table Architecture



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.47

Main Memory

47

<https://powcoder.com>

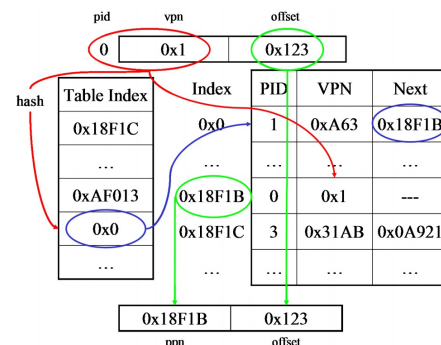


## Add WeChat powcoder Hashed Inverted Page Table

- Use hash table to limit the search to one — or at most a few — page-table entries

- TLB can accelerate access

- But how to implement shared memory?



<http://web.eecs.umich.edu/~akamil/teaching/sp04/040104.pdf>

EECS3221: Operating System Fundamentals

9.48

Main Memory

48





## Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
  - Total physical address space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.49

Main Memory

49

<https://powcoder.com>



## Add WeChat powcoder

### Swapping (Cont.)

- **Question:** does the swapped-out process need to swap back into the same physical addresses?
- Depends on address binding method
  - Plus consider pending I/O to-from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
  - Swapping normally disabled
  - Started if more than threshold amount of memory allocated
  - Disabled again once memory demand reduced below threshold

EECS3221: Operating System Fundamentals


9.50

Main Memory

50



## Schematic View of Swapping



Standard swapping involves moving entire processes between main memory and a backing store.

EECS3221: Operating System Fundamentals

9.51

Main Memory

51

<https://powcoder.com>



## Add WeChat powcoder

### Context Switch Time Including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
  - Swap out time of 2000 ms
  - Plus swap in of same sized process
  - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
  - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

EECS3221: Operating System Fundamentals

9.52

Main Memory

52



## Context Switch Time and Swapping (Cont.)

- Other constraints as well on swapping
  - Pending I/O – can't swap out as I/O would occur to wrong process
  - Or always transfer I/O to kernel space, then to I/O device
    - ▶ Known as **double buffering**, adds overhead
- Standard swapping not used in modern operating systems
  - But modified version common
    - ▶ Swap only when free memory extremely low

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.53

Main Memory

53

<https://powcoder.com>



## Swapping on Mobile Systems

- Not typically supported
  - Flash memory based
    - ▶ Small amount of space
    - ▶ Limited number of write cycles (SSD drives)
    - ▶ Poor throughput between flash memory and CPU on mobile platform
- Instead use other methods to free memory if low
  - iOS **asks** apps to voluntarily relinquish allocated memory
    - ▶ Read-only data thrown out and reloaded from flash if needed
    - ▶ Failure to free can result in termination
  - Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
  - Both OSes support paging as discussed below

EECS3221: Operating System Fundamentals


9.54

Main Memory

54



## Swapping with Paging



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.55

Main Memory

55

<https://powcoder.com>



## Add WeChat powcoder

Example: The Intel 32 and 64-bit Architectures

Dominant industry chips

Pentium CPUs are 32-bit and called IA-32 architecture

Current Intel CPUs are 64-bit and called IA-64 architecture

Many variations in the chips, cover the main ideas here

EECS3221: Operating System Fundamentals

9.56

Main Memory

56



## Example: The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging
  - Each segment can be 4GB
  - Up to 16K segments per process
  - The logical address space of a process is divided into two partitions.
    - ▶ First partition of up to 8K segments are private to process (kept in **local descriptor table (LDT)**)
    - ▶ Second partition of up to 8K segments shared among all processes (kept in **global descriptor table (GDT)**)

Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.57

Main Memory

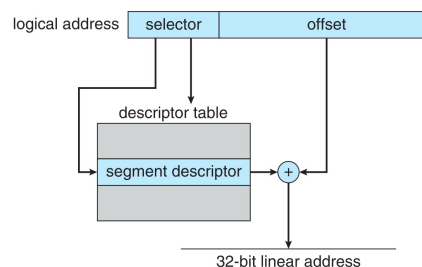

57

<https://powcoder.com>



## Add WeChat powcoder

### Intel IA-32 Segmentation



EECS3221: Operating System Fundamentals

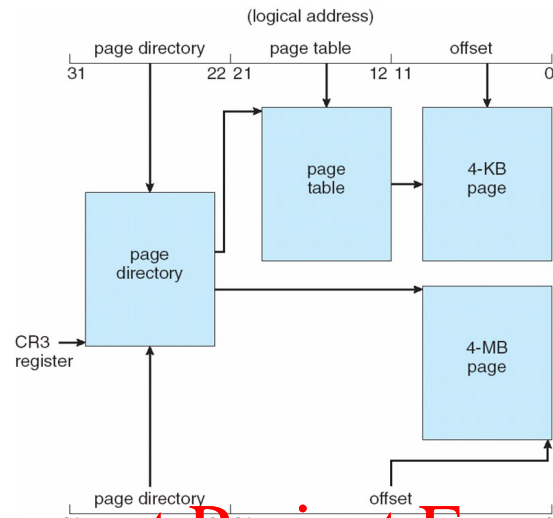
9.58

Main Memory

58



## Intel IA-32 Paging Architecture



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.59

Main Memory


59

<https://powcoder.com>



## Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space
  - Paging went to a 3-level scheme
  - Top two bits refer to a **page directory pointer table**
  - Page-directory and page-table entries moved to **64-bits** in size
  - Net effect is increasing address space to 36 bits – **64GB of physical memory**

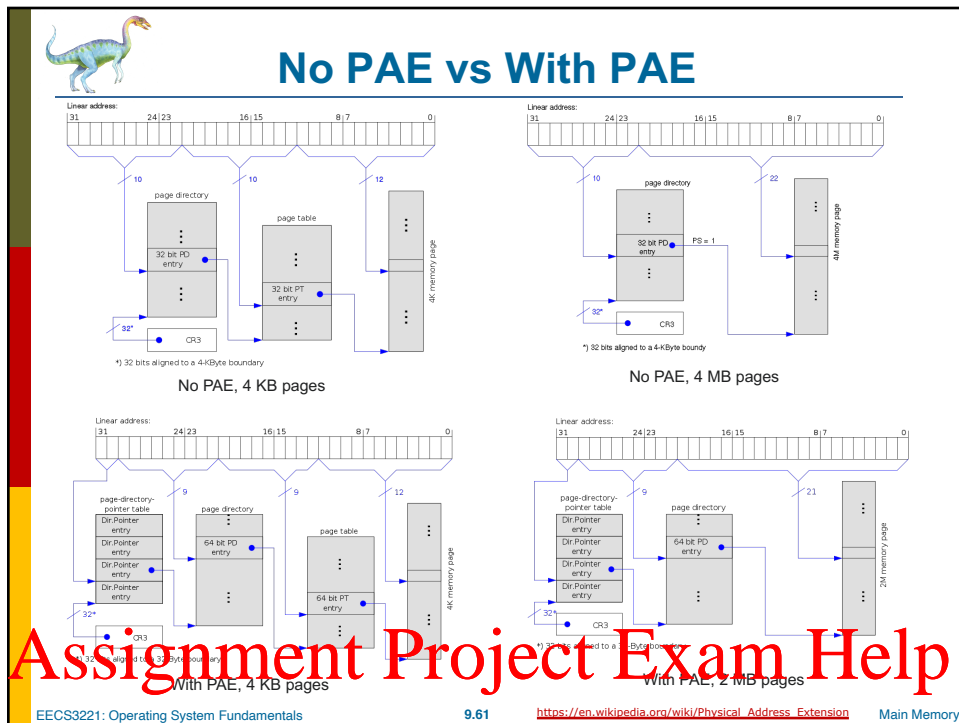


EECS3221: Operating System Fundamentals

9.60

Main Memory

60



61

<https://powcoder.com>

**Add WeChat powcoder**

Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
  - Page sizes of 4 KB, 2 MB, 1 GB
  - Four levels of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits

Common prefix			
Name	Symbol	Decimal	Binary
		SI	JEDEC
kilobyte	KB/kB	$10^3$	$2^{10}$
megabyte	MB	$10^6$	$2^{20}$
gigabyte	GB	$10^9$	$2^{30}$
terabyte	TB	$10^{12}$	$2^{40}$
petabyte	PB	$10^{15}$	$2^{50}$
exabyte	EB	$10^{18}$	$2^{60}$
zettabyte	ZB	$10^{21}$	$2^{70}$
yottabyte	YB	$10^{24}$	$2^{80}$

63 48 47 39 38 30 29 21 20 12 11 0

unused page map level 4 page directory pointer table page directory page table offset

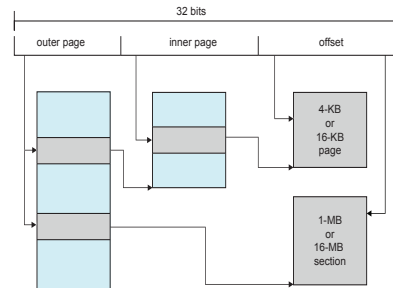
EECS3221: Operating System Fundamentals 9.62 Main Memory

62



## Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level for smaller pages



Assignment Project Exam Help

EECS3221: Operating System Fundamentals

9.63


Main Memory

63

<https://powcoder.com>

Add WeChat powcoder

## End of Chapter 9



64